Chemical Identity Lost in Regulation: A Study of Semantic Interoperability in European Chemical Substance Data

Abstract. Chemical substances appear across European environmental regulations, but regulatory datasets often represent them using heterogeneous identifiers and semantically mixed substance categories. This creates a major interoperability gap, as it remains unclear to what extent regulatory substance records can be linked to defined chemical structures across datasets. We analysed 18 ECHA and related datasets by mapping entries to InChIKeys where possible, classifying records by linkability, testing the consistency of CAS-number mappings, evaluating embedding-based semantic matching, and estimating the effort required for manual reconciliation and analysed whether commonly used identifiers or semantic methods can compensate for this mismatch. This paper presents a case study of semantic interoperability in European chemical substance data and a knowledge-graph-based analysis of how regulatory substance identity is represented across regulatory datasets. We found that only 62.5% of entries could be directly linked to a chemical structure, that CAS numbers do not function as reliably unique identifiers, that embedding-based approaches fail to recover chemical identity for non-structure substances, and that retrospective harmonisation would be infeasible at scale. These results show that interoperable and FAIR chemical regulation requires structure-based identifiers, formal ontologies, and machine-readable data models embedded at the level of regulatory design. As resource nonetheless, we include the scripts to generate an integrated knowledge graph from 18 regulatory sources.

Keywords. chemical substance identification, semantic interoperability, regulatory interoperability, REACH, InChIKey, cheminformatics, FAIR data, knowledge graph, European environmental legislation, Interoperable Europe

1. Introduction

Chemical substances are central entities in environmental, health, and product regulation. Across European legislation, substances are used to define obligations, assess risks, and communicate hazards in frameworks such as the Registration, Evaluation, Authorisation and Restriction of Chemicals (REACH) [1], the Classification, Labelling and Packaging of chemicals (CLP) [2], and sector-specific regulations including pesticides and water quality directives. Effective governance of these substances increasingly depends on the ability to integrate and reuse data across regulatory domains.

In chemistry and cheminformatics, chemical identity is well-defined and operationalised through algorithmically defined, structure-based, representations. Molecular structures can be uniquely encoded using standards such as SMILES notations [3], IUPAC names [4], the International Chemical Identifier (InChI) [5] and its hashed representation, the InChIKey, enabling unambiguous identification, comparison, and link-

ing of substances across databases [6]. Large-scale infrastructures such as PubChem [7] and ChEBI [8] build on these identifiers to provide machine-readable, ontology-based knowledge graphs that support interoperability and automated reasoning.

In contrast, regulatory data systems adopt a fundamentally different approach. Substances are typically identified using heterogeneous and often ambiguous identifiers, including Chemical Abstracts Service (CAS) registry numbers [9,10], European Community numbers (EC numbers) [11], textual names, and informal group descriptions. Moreover, the concept of a *substance* in regulation is not uniform: it may refer to a single molecule, a mixture of unknown or variable composition (UVCB) [12], a group of related compounds (e.g., PFAS), or even an analytical parameter defined as a function over multiple substances. As illustrated in Figure 1, these entities correspond to fundamentally different semantic types.

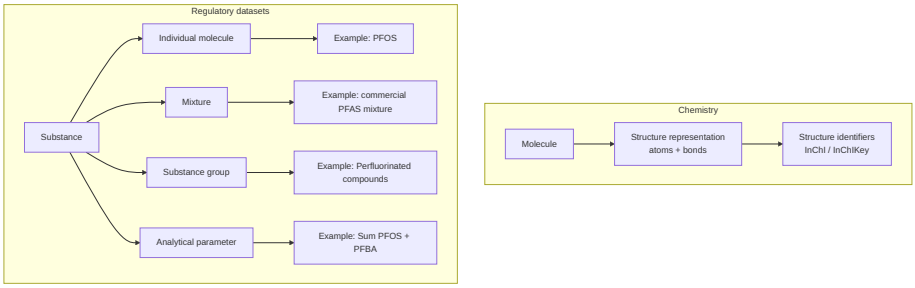

Figure 1. Conceptual model of the semantic mismatch. In chemistry, a substance typically corresponds to a molecule defined by its molecular structure and uniquely identified using structure-derived identifiers such as InChI or InChIKey. In regulatory datasets, however, the concept of a "substance" often encompasses heterogeneous entities including individual molecules, mixtures, substance groups, and analytical sum parameters. These entities correspond to fundamentally different semantic types, which complicates interoperability and automated reasoning across regulatory datasets. This conceptual mismatch is reflected empirically in Section 4, where a large fraction of regulatory entries cannot be linked to a defined molecular structure.

As a result, integrating substance information across regulatory frameworks becomes a complex and largely manual task, hindering policy enforcement for environmental administrations such as the Flemish Department of Environment and Spatial Development, and increasing uncertainty for industrial players in environmental management, investment, and innovation.

This paper investigates the extent and consequences of this interoperability problem using a comprehensive analysis of 18 regulatory datasets from the European Chemicals Agency (ECHA) and related sources. We focus on three main research questions:

- To what extent can regulatory substances be linked to unambiguous chemical structures?
- How consistent and reliable are commonly used identifiers such as CAS numbers for cross-dataset linking?
- Can semantic or data-driven methods compensate for the absence of structure-based identifiers?

Based on this analysis, we make the following contributions:

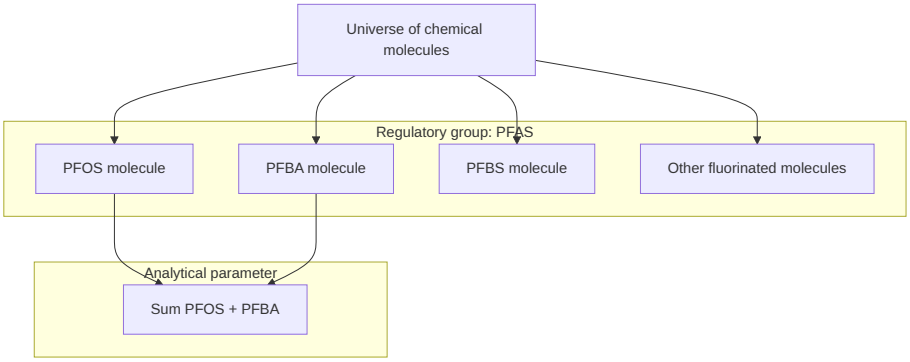

Figure 2. Set-theoretic view of regulatory substances. The lack of explicit set-based semantics leads to unclear relationships between entities and contributes to combinatorial complexity.

1. We quantify the heterogeneity of regulatory substance definitions and introduce a taxonomy of linkability reflecting their degree of chemical specificity.
2. We demonstrate the limitations of CAS-based identification and show that a substantial fraction of regulatory substances cannot be resolved to a defined structure.
3. We evaluate semantic embedding approaches and show that they fail to recover chemical identity for non-structure substances.
4. We estimate the scale of the manual effort required for retrospective harmonisation, demonstrating its infeasibility.
5. We show that structure-based identifiers combined with ontology-driven approaches enable scalable integration through knowledge graphs.

The remainder of this paper is structured as follows. Section 3 describes the data sources and analytical methods. Section 4 presents the empirical findings. Section 5 discusses the implications for regulatory data infrastructures and semantic interoperability. Section 6 concludes with recommendations for future work.

2. Related Work

Interoperability of scientific data has been extensively studied in the context of FAIR principles [13] and knowledge graph infrastructures. In the life sciences, large-scale resources such as PubChem [7] and the Chemical Entities of Biological Interest (ChEBI) ontology [8] demonstrate how chemical entities can be represented using structure-based identifiers, ontology-driven classifications, and machine-readable knowledge graphs. Extensive work on rule-based chemical classification has been undertaken by ClassyFire [14]. In this framework, a structure-based taxonomy (ChemOnt) has been established, where each class is characterised by precise and unambiguous formal rules. These infrastructures enable consistent integration, querying, and reuse of chemical information across domains.

Within the Semantic Web community, numerous approaches have been proposed for data integration based on linked data principles, ontology alignment, and entity res-

olution. Techniques such as identifier mapping, schema matching, and, more recently, embedding-based similarity methods [15] have been applied to address heterogeneity across datasets. However, these approaches typically assume that entities correspond to well-defined and comparable concepts.

In cheminformatics and regulatory science, considerable effort has focused on handling substances without precise structural definitions, such as UVCBs. Methods like those of Kutsarova et al. [12] use generic structural encodings (e.g., G SMILES) and statistical sampling to generate representative constituent structures, reducing combinatorial complexity while preserving structural diversity.

These approaches operate at the level of individual substances and do not address broader semantic inconsistencies across regulatory datasets. Many entries still lack structure-resolvable identifiers or rely on heterogeneous definitions, limiting interoperability between data sources.

In the regulatory domain, previous work has highlighted challenges related to data fragmentation, heterogeneity of identifiers, and limited machine-readability of legal data [16]. In particular, chemical regulation introduces additional complexity due to the use of substance groups, mixtures, and analytical parameters that do not correspond to single chemical structures. While prior studies have focused on improving data accessibility or applying FAIR principles to regulatory datasets, the fundamental mismatch between chemical identity and regulatory substance definitions has received limited attention.

3. Methods

Our analysis comprises the following steps:

Data Integration Substance records were compiled from a comprehensive set of 18 European regulatory and monitoring sources and one Flemish administrative list. These include ECHA obligation lists (Candidate List [17], Authorisation List [18], Restriction List [19], POPs List [20], EU Positive List [21], and the CLH Harmonised Classification and Labelling List [22]), ECHA activity lists (Restriction Process [23], SVHC Identification [24], Authorisation Process [25], Dossier Evaluation [26], CLH Process [27], Substance Evaluation [28], POPs Process [29], PBT Assessment [30], and ED Assessment [31]), as well as REACH registration data [32].

Additional sources comprise the EU Pesticides Database (Active Substances) [33], the OSPAR List of Chemicals for Priority Action [34], and a Flemish administrative substance list developed in the context of water quality modelling [35]. This Flemish `summationsubstances` list reflects the heterogeneity of *substance names* in Flemish environmental permits and monitoring. For each entry, available identifiers (CAS number, EC number, and substance name) were collected.

Where possible, substances were linked to structure-based identifiers (InChIKey) via PubChem [7]. The resulting dataset establishes a unified mapping between regulatory entries, their identifiers, and their originating source lists.

Chemical Identity and Linkability Each substance entry was classified into one of several types reflecting its degree of chemical specificity: structure-defined molecule, substance group, mixture/UVCB, analytical parameter, or administrative entry.

Based on this classification, a linkability taxonomy was defined, ranging from fully structure-resolvable entries (InChIKey) to non-identifiable entries. This taxonomy expresses the extent to which regulatory substances can be unambiguously linked across datasets.

Identifier Consistency Analysis For entries with both CAS numbers and InChIKeys, bidirectional mappings were analysed to assess identifier consistency. This analysis evaluates whether CAS numbers uniquely map to structures and vice versa.

Structure Coverage and Cross-List Overlap For each regulatory source, the proportion of structure-defined substances was calculated. Overlap between sources was analysed using structure-based matching (InChIKey) and set similarity measures.

Analysis of Non-Structure Substances Substances without structure identifiers were analysed using text-based methods, including embedding-based clustering [36,37,38] and similarity matching to chemical taxonomies (ChemOnt), to assess whether semantic techniques can recover meaningful structure or classification [39,14]. Each substance name and ChemOnt class label was encoded using the `all-MiniLM-L6-v2` sentence-transformer model [15,40,41], yielding a 384-dimensional dense vector representation.

Embedding Model Validation Using ChemOnt Ground Truth To evaluate the ability of text-based embeddings to recover chemical classifications from substance names, an internal validation was performed using structure-defined substances as ground truth. Substances with an unambiguous InChIKey were linked to the ChemOnt taxonomy through ClassyFire classification [14], providing reference assignments at the levels of direct parent, subclass, class, superclass, and kingdom.

Substance names and ChemOnt class labels were encoded using pre-trained transformer models (`all-MiniLM-L6-v2`, `all-mpnet-base-v2`, `allenai/scibert_scivocab_uncased`, and `dmis-lab/biobert-v1.1`) within the *sentence-transformers* framework [15,40,41]. Cosine similarity was used to rank all ChemOnt classes for each substance name.

Performance was assessed using standard information retrieval metrics, including Hit@k ($k = 1, 3, 5, 10$) and Mean Reciprocal Rank (MRR) [39]. To account for the hierarchical structure of ChemOnt, Hit@1 was additionally computed for each taxonomy level. A weighted lenient MRR was introduced to summarise performance across levels, assigning ordinal weights of 1.00 (direct parent), 0.75 (subclass), 0.50 (class), 0.25 (superclass), and 0.10 (kingdom), thereby favouring more specific classifications.

This validation provides a principled comparison of embedding models and demonstrates the extent to which semantic representations of substance names approximate structure-based chemical taxonomy assignments.

Knowledge Graph Construction All data were integrated into an RDF knowledge graph using SKOS [42] for taxonomy representation. Structure-defined substances were linked to ChemOnt classes and to ChEBI entities via InChIKey, enabling integration of biological hazard annotations [43].

Workload Estimation For non-structure substances, interoperability requires manual assessment of pairwise relationships (equivalence, subset, overlap, disjoint). The required effort scales quadratically with the number of such entries as $\binom{N}{2} = N(N-1)/2$, where N is the number of non-structure entries under consideration.

Priority Scoring To illustrate the analytical value of the integrated knowledge graph (C5), a composite priority score was computed by combining regulatory list membership with biological role annotations from ChEBI [8], following the ToxPi framework [44]. This score is illustrative only and does not constitute a formal risk assessment.

4. Results

Overview of the integrated knowledge graph We constructed an integrated knowledge graph to analyse the representation and interoperability of regulatory substances across datasets. Figure 3 provides an overview of the resulting graph structure and its main semantic layers.

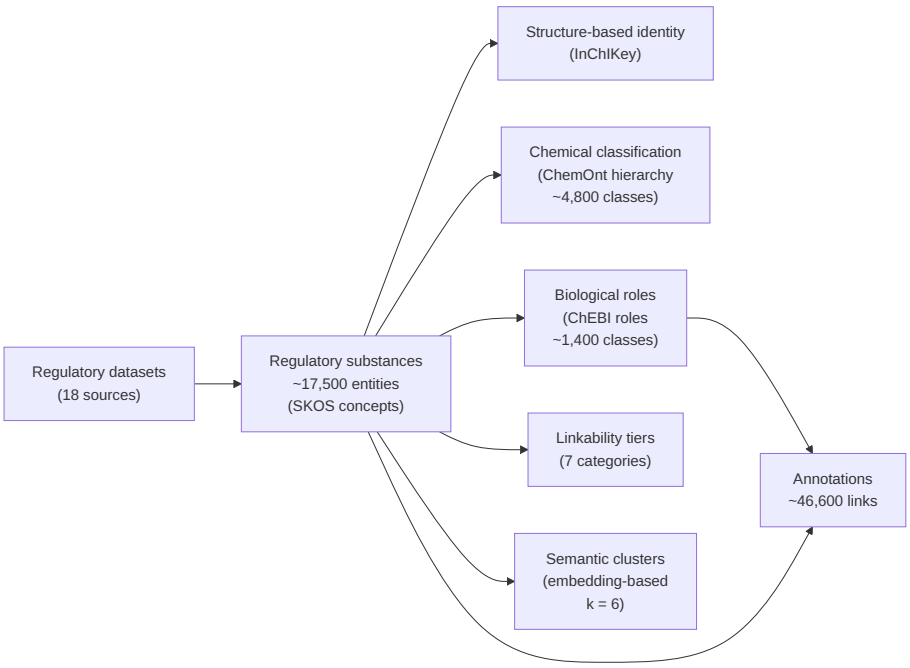

Figure 3. Overview of the integrated knowledge graph constructed from 18 regulatory datasets. Regulatory substances (17,500 entities) are represented as SKOS concepts and form the core layer of the graph. Where possible, substances are linked to structure-based identifiers (InChIKey), enabling alignment with chemical classification systems (ChemOnt) and biological roles (ChEBI). In addition, substances are organised into linkability tiers reflecting their degree of chemical specificity, and into semantic clusters derived from embedding-based similarity. The graph is further enriched with annotations (46,600 links) connecting substances to regulatory sources and functional roles. The figure highlights the multiple semantic layers required to represent and integrate heterogeneous regulatory substance definitions.

Dataset Overview (C1) Integration of 18 regulatory sources yielded 41,813 records. Of these, 63% could be linked to a defined molecular structure via InChIKey.

Heterogeneity of Substance Definitions A significant share of regulatory entries (37%) does not correspond to a defined chemical structure, consistent with Kutsarova et al. [12]. These include mixtures (UVCBs), substance groups, analytical parameters, or entries without usable identifiers, reflecting differences in how "substances" are defined across regulatory contexts.

Heterogeneity varies widely between regulatory lists: some are mostly structure-based, while others rely heavily on group or mixture definitions, limiting interoperability (Figure 2).

Applying the seven-tier linkability taxonomy (Section 3) shows that 62.5% of entries (20,909) are directly linkable via InChIKey (tier 1). Only a small fraction consists of partially resolvable groups (tiers 2–3; <0.2%).

The remaining 37.5% cannot be directly linked to a structure: 11.6% have unresolved CAS numbers (tier 4), 10.8% are mixtures or UVCBs (tier 5), <0.1% are administrative entries (tier 6), and 14.9% lack any usable identifier (tier 7).

More than one-third of regulatory entries cannot be mapped to a unique chemical structure, and even CAS numbers do not guarantee unambiguous identification. Consequently, structure-based integration applies only to part of the data, requiring alternative approaches for the remainder.

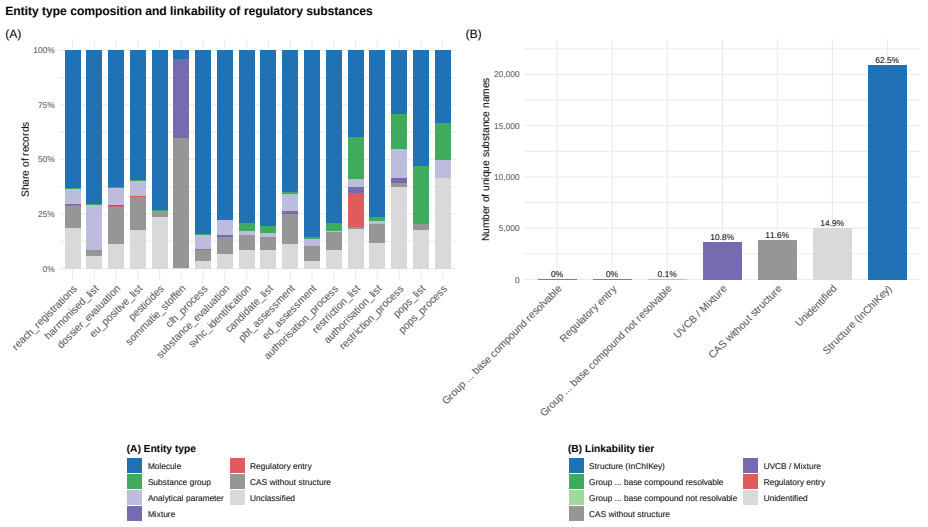

Figure 4. (A) Entity type composition across 18 regulatory sources. Each bar shows the proportional breakdown of substance entries into seven mutually exclusive types (Molecule, Substance group, Analytical parameter, Mixture, Regulatory entry, CAS without structure, Unclassified), illustrating that the mix of entity types differs substantially between regulatory instruments. (B) Linkability taxonomy applied to all 33,452 unique regulatory substance entries. A total of 20,909 entries (62.5%) can be directly linked to a chemical structure via InChIKey (tier 1), while 4,981 entries (14.9%) lack any usable identifier. Together, the panels show that structural linkability is unevenly distributed across sources and that more than one-third of all entries cannot be resolved to a unique chemical structure.

Limitations of CAS Numbers (C2) CAS numbers are not reliable unique identifiers. Bidirectional mapping analysis reveals that 195 InChIKeys are associated with more than

one CAS number (up to 7 per structure), while 4 CAS numbers resolve to two distinct InChIKeys. These inconsistencies arise from differences in stereochemistry, salt forms, and context-dependent registration practices. CAS-based cross-dataset linking therefore introduces ambiguity rather than resolving it, making CAS numbers insufficient as a backbone for interoperability.

Limited Interoperability Across Regulatory Lists Cross-list analysis of the 17,523 structure-resolved InChIKeys shows that 70.0% (12,270 entries) appear in only a single regulatory source, while just 30.0% (5,253 entries) are shared across two or more lists. This limited overlap reflects the domain-specific design of regulatory instruments rather than a shared chemical reference framework.

Failure of Semantic Approaches for Non-Structure Substances (C3) Embedding-based clustering of non-structure substance names yields semantically coherent groups, but does not resolve chemical identity. While clusters reflect lexical or contextual similarity between substance descriptions, they do not correspond to well-defined molecular entities (Figure 5A).

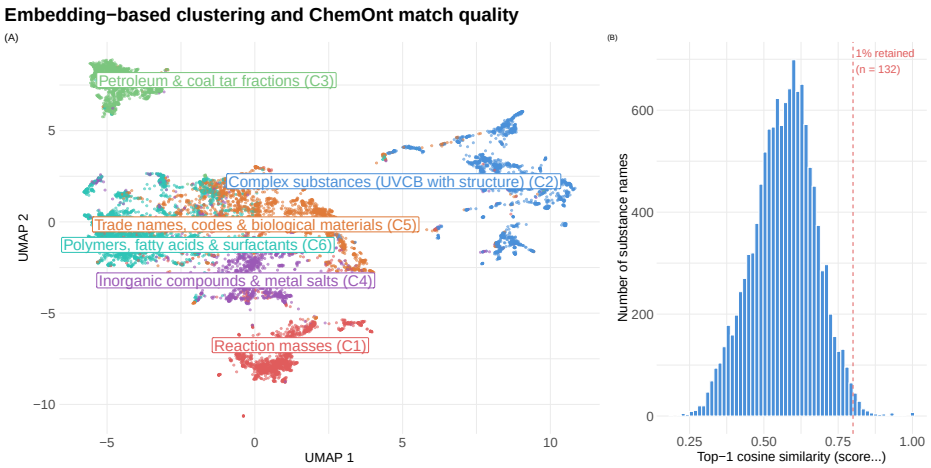

Figure 5. (A) UMAP projection of sentence embeddings for 12,499 non-structure regulatory substance entries, coloured by k-means cluster (k = 6). Each point represents a unique substance name; spatial proximity reflects semantic similarity in the 384-dimensional embedding space. (B) Distribution of top-1 cosine similarity scores [39] between substance name embeddings and ChemOnt class label embeddings across all 12,499 candidate substance-class pairs, shown here for `all-MiniLM-L6-v2`. The dashed line indicates the applied threshold ($s_1 > 0.80$); 132 entries (1.1%) are retained as matches.

To evaluate whether semantic similarity can recover chemical meaning, we attempted to map non-structure substances to ChemOnt taxonomy classes using cosine similarity between sentence embeddings of substance names and ChemOnt class labels. Four models were evaluated: `all-MiniLM-L6-v2` (miniLM), `all-mpnet-base-v2` (mpnet), `allenai/scibert\_scivocab\_uncased` (SciBERT), and `dmis-lab/biobert-v1.1` (BioBERT). The distribution of top-1 similarity scores is heavily skewed towards lower values (Figure 5B). Cosine similarity scores are not comparable across models: scores for the 200th-ranked match range from approximately 0.80 for miniLM

to 0.94 for BioBERT, reflecting differences in embedding geometry rather than differences in match quality. A fixed threshold such as $s_1 > 0.80$ is therefore model-dependent and arbitrary; it was applied here for the purpose of bounding a manual evaluation, not as a principled quality criterion.

Manual inspection of the top-200 matches per model revealed that the majority of retained assignments are chemically incorrect, regardless of model. Correct assignments were largely limited to cases where the substance name closely resembles the ChemOnt class label (e.g. *Phenols* → *Phenols*, *Alkyl mercury compounds* → *Alkyl mercury compounds*), indicating that high similarity scores reflect surface-form overlap rather than semantic understanding. Systematic errors were present across all models: miniLM conflates organic and inorganic species and confuses elements sharing a common prefix (e.g. *organostannic compounds* → *organoarsenic compounds*); mpnet consistently misassigns all nickel sulfide variants to *nitrile sulfides*, a token-level artefact of the chemical symbol Ni being read as the start of "nitrile"; SciBERT and BioBERT, pre-trained on biomedical literature, exhibit severe score inflation—a score of 0.93 on an incorrect match provides less information than a score of 0.85 on a correct one—and produce repeated false anchors to specific heterocyclic and lipid ChemOnt classes for structurally unrelated substances. Trade names and reaction products were assigned arbitrary classes with no verifiable basis across all models.

Score magnitude is not a reliable indicator of match correctness. Models with the highest score floors (SciBERT, BioBERT) performed worst on both manual inspection and quantitative validation, demonstrating that inflated cosine scores reflect poor calibration rather than genuine chemical understanding. Embedding-based methods cannot compensate for the absence of structure-based identifiers, and fail systematically for substance groups, mixtures, and analytically defined entries.

To validate this failure quantitatively, the four models were evaluated against a structure-derived ChemOnt ground truth comprising 28,204 ClassyFire-assigned substances (Table 1). `all-MiniLM-L6-v2` achieved the highest Hit@1 of 5.75% and MRR of 0.093 — even this best result implies the correct class is retrieved in fewer than 1 in 17 attempts. All other models scored below Hit@1 = 1%. Retrieval performance decreased further with increasing taxonomic abstraction: Hit@1 fell to 3.25% at the subclass level and below 1% for class, superclass, and kingdom. A hierarchy-aware weighted MRR of 0.099 was only marginally higher than the strict MRR, confirming that partial matches at higher levels of the ChemOnt hierarchy do not rescue overall classification quality.

Table 1. Performance of embedding models for ChemOnt classification.

Model	Hit@1	Hit@3	Hit@5	Hit@10	MRR
all-MiniLM-L6-v2	0.0575	0.0986	0.1234	0.1621	0.093
all-mpnet-base-v2	0.0093	0.0152	0.0184	0.0247	0.015
SciBERT	0.0040	0.0071	0.0102	0.0131	0.008
BioBERT	0.0047	0.0077	0.0099	0.0150	0.009

Scale of the Interoperability Problem (C4) For the more than 11,000 non-structure entries, full manual reconciliation would require on the order of hundreds of millions of pairwise comparisons. Under conservative assumptions, this corresponds to several thousand person-years of expert effort.

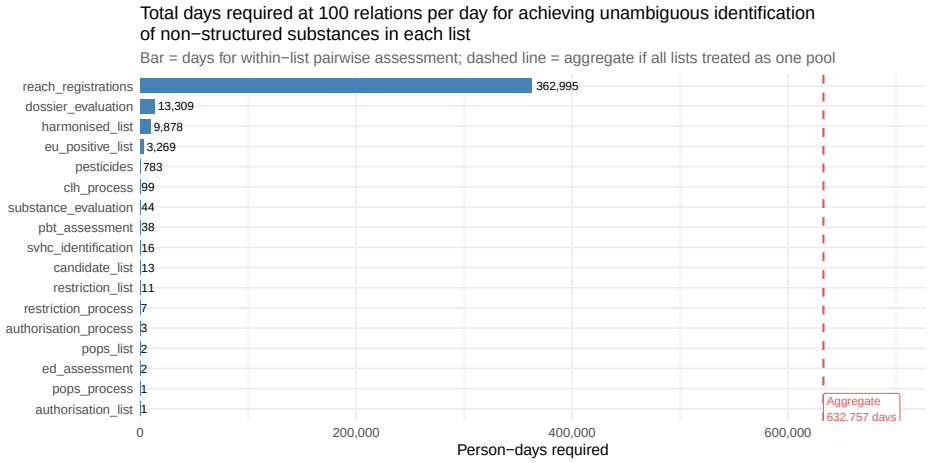

Figure 6. Estimated pairwise curation workload for 11,250 non-structure entries across the combined European regulatory dataset, broken down by source. The total corresponds to approximately 630,000 person–days (3,150 person-years) under an assumed throughput of 100 pair assessments per person-day. The quadratic growth of the workload illustrates that manual reconciliation is not scalable, reinforcing the need for structure-based and ontology-driven representations at the source.

Figure 6 quantifies this workload across all regulatory sources, illustrating both the absolute scale and the quadratic growth that makes each additional non-structure entry disproportionately costly. Even moderate increases in the number of unresolved entries lead to a rapidly escalating curation burden.

Manual pairwise reconciliation at this scale is not operationally feasible. This result highlights a fundamental constraint: interoperability cannot be achieved through retrospective harmonisation of heterogeneous substance definitions.

Instead, the only scalable strategy is to address the problem at its source. By requiring structure-resolvable identifiers and ontology-based representations in regulatory definitions wherever chemically applicable, the number of entries requiring manual reconciliation can be structurally reduced rather than managed post hoc.

Knowledge Graph Integration and Priority Scoring (C5) Structure-defined substances can be successfully integrated across regulatory lists, chemical taxonomies (ChemOnt), and biological annotations (ChEBI) using InChIKey as a linking key. This enables cross-domain queries that are infeasible on fragmented regulatory datasets: for example, identifying substances that appear on three or more ECHA lists and carry a documented biological role in ChEBI, or retrieving all members of a given ChemOnt class that are subject to authorisation requirements. As an illustration of this analytical potential, a composite priority score combining regulatory list membership with ChEBI biological role annotations was computed following the ToxPi framework [44]. Figure 7 plots regulatory versus biological scores for all scored substances; bisphenol A and dibutyl phthalate exemplify substances with both broad regulatory presence and multiple documented biological roles. This score is illustrative and does not constitute a formal risk assessment.

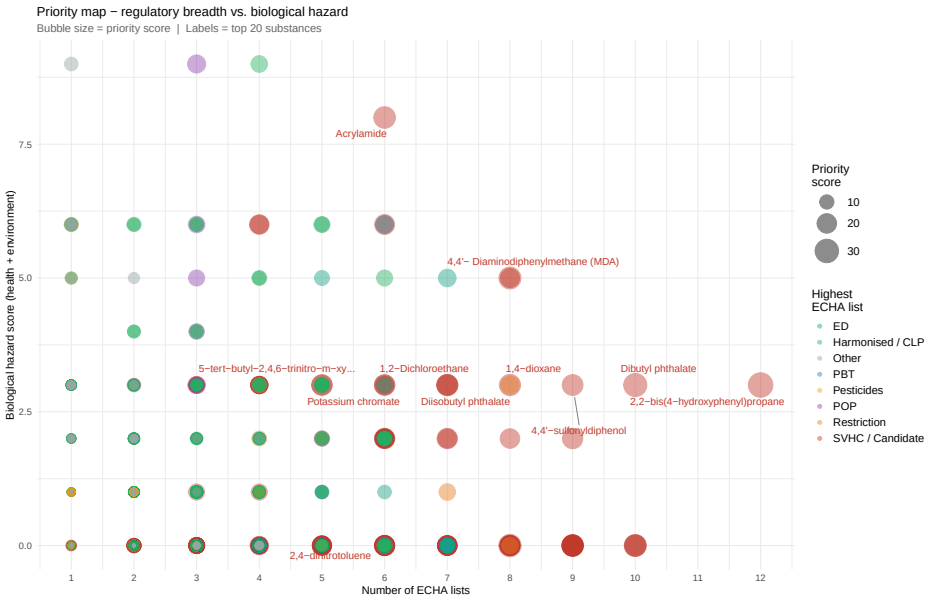

Figure 7. Bubble chart of composite priority space: regulatory list score versus biological hazard score for all scored substances. Bubble size is proportional to the number of regulatory source memberships. Substances in the upper-right quadrant combine broad regulatory coverage with documented biological hazard evidence.

5. Implications

Structural linkability and its limits (RQ1) The results demonstrate that interoperability of regulatory substance data is fundamentally constrained by the absence of a shared and formalised notion of chemical identity. While structure-based identifiers such as InChIKey enable unambiguous linking across datasets, a substantial fraction of regulatory entries cannot be expressed in this form. Instead, regulatory "substances" frequently represent heterogeneous entities, including mixtures, substance groups, and analytical parameters.

This heterogeneity is not merely a data quality issue but reflects a structural mismatch between chemical reality and regulatory abstraction. As a consequence, semantic relationships between regulatory entries, such as equivalence, subset relations or partial overlap, remain implicit and undefined. Without explicit formalisation of these relationships, automated integration across regulatory datasets is inherently limited.

Unreliability of CAS-based identification (RQ2) The analysis confirms that commonly used identifiers in regulatory contexts, particularly CAS numbers, do not provide a reliable foundation for interoperability. Although widely adopted, they exhibit both one-to-many and many-to-one mappings with chemical structures, reflecting differences in stereochemistry, salt forms, and context-dependent definitions.

This finding challenges the widespread assumption in regulatory data integration that such identifiers can serve as globally unique references. In practice, relying on them introduces ambiguity rather than resolving it, undermining their suitability as a primary reference mechanism for cross-dataset linking.

Failure of semantic compensation (RQ3) This study shows that modern embedding-based methods cannot compensate for the absence of structure-based identifiers in regulatory substance data. While they can group substance names into loosely coherent semantic clusters, they fail to recover true chemical identity or reliably map entries to ChemOnt classes, with only a negligible number of correct assignments.

These limitations are inherent to the data rather than the models: textual descriptions of mixtures, groups, and analytical entries lack the information required to infer molecular structure. As a result, semantic similarity largely reflects lexical overlap rather than chemical relatedness.

Sentence embeddings cannot therefore substitute for structure-resolvable identifiers. Reliable chemical classification and interoperability therefore require explicit structural representation, such as InChIKeys, at the source.

The combinatorial barrier to retrospective harmonisation The scale of the interoperability problem becomes evident when considering the effort required for manual reconciliation. For non-structure substances, establishing semantic relationships requires pairwise comparison, leading to quadratic growth in workload. The estimated effort of several thousand person-years demonstrates that retrospective correction is not a viable strategy at scale, interoperability must be built into regulatory data models by design.

Towards ontology-driven regulatory data infrastructures In contrast to regulatory datasets, cheminformatics infrastructures such as PubChem and ChEBI demonstrate that large-scale interoperability is achievable when chemical identity is explicitly represented and linked to formal ontologies. Structure-based identifiers, combined with hierarchical classifications and machine-readable knowledge graphs, enable consistent integration across domains.

The priority scoring example illustrates that once chemical identity is formalised and linked across datasets, higher-level aggregations combining regulatory and biological information become straightforward to compute.

The results of this study show that similar principles can be applied to regulatory data. Structure-defined substances can be successfully linked across datasets and enriched with chemical and biological annotations when represented using InChIKey and ontology-based classifications.

Implications for FAIR chemical data in Europe The findings have direct implications for the implementation of FAIR (Findable, Accessible, Interoperable, Reusable) principles in European chemical regulation. Current regulatory datasets fall short of interoperability primarily because they lack globally unique, persistent, and machine-readable identifiers grounded in chemical structure.

Achieving FAIR chemical substance data requires a shift from document-centric and identifier-fragmented approaches towards linked data infrastructures in which substances are represented as resolvable URIs, grounded in structure-based identifiers and connected through formal ontologies. This shift is not purely technical: it requires alignment between data modelling practices and the legal definitions of substances in regulatory frameworks.

Limitations and scope This study focuses on datasets derived from ECHA and related European regulatory sources, and on the use of InChIKey as the primary structure-based identifier. While the results are consistent across the analysed datasets, further work is

needed to assess the generality of these findings across other regulatory domains and jurisdictions.

Furthermore, not all regulatory concepts can or should be reduced to individual molecular structures. Mixtures, UVCBs, and analytical parameters are legitimate and necessary constructs in regulation. However, their current representation lacks formal semantics, limiting their integration with structure-based systems. Future work should explore hybrid models that explicitly represent these entities as sets, classes, or functions over well-defined chemical structures.

From data integration to regulatory design The interoperability problem is not primarily a technical integration challenge, but a consequence of how substances are defined and represented in regulation. Efforts to improve interoperability must therefore extend beyond data engineering to include the design of regulatory concepts and data standards.

Embedding structure-based identifiers, ontology references, and machine-readable semantics directly into regulatory processes would significantly reduce ambiguity and enable scalable integration. Without such changes, the gap between regulatory data and scientific data infrastructures will continue to widen, limiting the effectiveness of data-driven environmental governance.

In practice, this implies that regulatory data infrastructures should: (i) include structure-based identifiers (e.g. InChIKey) wherever applicable; (ii) represent substance groups and mixtures using explicit semantic models (e.g. set- or class-based representations); (iii) publish substance data as linked data with persistent URIs.

6. Conclusion

This study shows that interoperability of chemical substance data in European regulatory frameworks is fundamentally limited by the lack of a shared and formalised notion of chemical identity. Analysis of 18 regulatory datasets reveals that a substantial fraction of substances cannot be unambiguously linked to a molecular structure, and that commonly used identifiers such as CAS numbers are insufficient for reliable cross-dataset integration.

The three research questions posed in the introduction are answered as follows. RQ1: 37.5% of regulatory entries cannot be directly linked to a defined chemical structure; 62.5% resolve via InChIKey, while the remaining entries represent mixtures, substance groups, unresolved CAS numbers, or entries without usable identifiers. RQ2: CAS numbers prove unreliable as unique identifiers; bidirectional mapping shows that 195 InChIKeys are linked to more than one CAS number and 4 CAS numbers resolve to two distinct InChIKeys, with many-to-one mappings arising from differences in stereochemistry, salt forms, and registration context. RQ3: semantic and embedding-based approaches cannot compensate for the absence of structure-based identifiers; the best-performing model achieves a ChemOnt Hit@1 of only 5.75%, and manual inspection reveals systematic errors across all four evaluated models, with high cosine similarity scores reflecting surface-form overlap rather than chemical understanding.

These limitations directly affect semantic interoperability. Many regulatory definitions lack explicit structure-based identifiers and formal semantics, leaving relationships such as equivalence, inclusion, or overlap implicit. This is particularly evident for substance groups and analytical parameters, which often represent sets or transformations

of underlying chemical entities. Without an explicit formalisation of these relationships, interoperability and automated reasoning remain severely constrained.

The problem is amplified by the heterogeneity and scale of regulatory data, where multiple partially overlapping substance lists are defined in different legal and domain-specific contexts. In this setting, both manual curation and data-driven approaches face intrinsic scalability limits, making retrospective harmonisation insufficient.

Our results indicate that interoperability cannot be achieved through post hoc technical solutions alone, including semantic embedding approaches. Instead, it requires changes at the level of data modelling and regulation. The use of structure-based identifiers (e.g. InChIKeys), combined with formal ontologies and linked data principles, is necessary to enable consistent integration when applied at source.

We therefore argue that achieving FAIR and interoperable chemical data requires a shift towards embedding these principles directly into regulatory data standards and legislative processes. Without such changes, fragmentation between regulatory and scientific data infrastructures will persist.

Future work should focus on formal representations for mixtures and substance groups, and on translating these into operational standards for regulatory data exchange.

Data Availability

All data, code, and materials supporting this study are publicly available. The complete dataset, including the integrated knowledge graph, analysis scripts, and figures, is archived on Zenodo at `https://doi.org/10.5281/zenodo.xxxxxxx`.

Supplementary Materials

The integrated knowledge graph (RDF/Turtle), processed datasets from 18 regulatory sources, analysis scripts, and all figures are available via the Zenodo archive associated with this publication.

Declaration on the Use of Generative AI

During the preparation of this work, the author used Claude Code (Anthropic, model: *claude-sonnet-4-6*) for the following purposes: (1) writing and debugging analysis code in R, Bash, and Python; (2) assisting with the configuration of a Linux environment, including the NVIDIA/CUDA software stack installed on an external SSD to enable local execution of transformer-based embedding models; and (3) translating draft text originally written in Dutch into British English. After using this tool, the author reviewed and edited all content as needed and takes full responsibility for the publication's content.

References

[1] European Parliament, Council of the European Union. Regulation (EC) No 1907/2006 concerning the Registration, Evaluation, Authorisation and Restriction of Chemicals (REACH). Official Journal of the European Union; 2006. L 396. Available from: http://data.europa.eu/eli/reg/2006/1907/2022-03-01.

[2] European Parliament, Council of the European Union. Regulation (EC) No 1272/2008 on classification, labelling and packaging of substances and mixtures (CLP). Official Journal of the European Union; 2008. L 353. Available from: http://data.europa.eu/eli/reg/2008/1272/oj.

[3] Weininger D. SMILES, a chemical language and information system. 1. Introduction to methodology and encoding rules. Journal of Chemical Information and Computer Sciences. 1988 Feb;28(1):31-6. Available from: https://doi.org/10.1021/ci00057a005.

[4] Koppenol WH, Corish J, García-Martínez J, Meija J, Reedijk J. How to name new chemical elements (IUPAC Recommendations 2016). Pure and Applied Chemistry. 2016 Apr;88(4):401-5.

[5] Heller SR, McNaught A, Pletnev I, Stein S, Tchekhovskoi D. InChI, the IUPAC International Chemical Identifier. Journal of Cheminformatics. 2015 May;7(1):23. Available from: https://doi.org/10.1186/s13321-015-0068-4.

[6] Akhondi SA, Kors JA, Muresan S. Consistency of systematic chemical identifiers within and between small-molecule databases. Journal of Cheminformatics. 2012 Dec;4(1):35. Available from: https://doi.org/10.1186/1758-2946-4-35.

[7] Kim S, Thiessen PA, Bolton EE, Bryant SH. PUG-SOAP and PUG-REST: web services for programmatic access to chemical information in PubChem. Nucleic Acids Research. 2015 Jul;43(W1):W605-11.

[8] Degtyarenko K, de Matos P, Ennis M, Hastings J, Zbinden M, McNaught A, et al. ChEBI: a database and ontology for chemical entities of biological interest. Nucleic Acids Research. 2008 Jan;36(suppl_1):D344-50. Available from: https://doi.org/10.1093/nar/gkm791.

[9] Chemical Abstracts Service. CAS Registry; 2024. Accessed: 2026-04-08. https://www.cas.org.

[10] Dittmar PG, Stobaugh RE, Watson CE. The Chemical Abstracts Service Chemical Registry System. I. General Design. Journal of Chemical Information and Computer Sciences. 1976 May;16(2):111-21. Available from: https://doi.org/10.1021/ci60006a016.

[11] European Chemicals Agency (ECHA). About EC and List Numbers; n.d. Accessed April 8, 2026. https://echa.europa.eu/en/about-ec-and-list-numbers.

[12] Kutsarova SS, Yordanova DG, Karakolev YH, Stoeva S, Comber M, Hughes CB, et al. UVCB substances II: Development of an endpoint-nonspecific procedure for selection of computationally generated representative constituents. Environmental Toxicology and Chemistry. 2019 Mar;38(3):682-94. Available from: https://doi.org/10.1002/etc.4358.

[13] Wilkinson MD, Dumontier M, Aalbersberg IJ, Appleton G, Axton M, Baak A, et al. The FAIR Guiding Principles for scientific data management and stewardship. Scientific Data. 2016 Mar;3(1):160018. Available from: https://doi.org/10.1038/sdata.2016.18.

[14] Djoumbou Feunang Y, Eisner R, Knox C, Chepelev L, Hastings J, Owen G, et al. ClassyFire: automated chemical classification with a comprehensive, computable taxonomy. Journal of Cheminformatics. 2016 Nov;8(1):61. Available from: https://doi.org/10.1186/s13321-016-0174-y.

[15] Reimers N, Gurevych I. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing. Association for Computational Linguistics; 2019. p. 3982-92.

[16] Filtz E, Kirrane S, Polleres A. The linked legal data landscape: linking legal data across different countries. Artificial Intelligence and Law. 2021 Dec;29(4):485-539. Available from: https://doi.org/10.1007/s10506-021-09282-8.

[17] European Chemicals Agency (ECHA). Candidate List of Substances of Very High Concern; 2025. https://chem.echa.europa.eu/obligation-lists/candidateList.

[18] European Chemicals Agency (ECHA). Authorisation List; 2025. https://chem.echa.europa.eu/obligation-lists/authorisationList.

[19] European Chemicals Agency (ECHA). Restriction List; 2025. https://chem.echa.europa.eu/obligation-lists/restrictionList.

[20] European Chemicals Agency (ECHA). Persistent Organic Pollutants (POPs) List; 2025. https://chem.echa.europa.eu/obligation-lists/popsList.

[21] European Chemicals Agency (ECHA). EU Positive List; 2025. https://chem.echa.europa.eu/obligation-lists/euPositiveList.
[22] European Chemicals Agency (ECHA). Harmonised Classification and Labelling List (CLH); 2025. https://chem.echa.europa.eu/obligation-lists/clhList.
[23] European Chemicals Agency (ECHA). Restriction Process List; 2025. https://chem.echa.europa.eu/activity-lists/restrictionProcess.
[24] European Chemicals Agency (ECHA). SVHC Identification List; 2025. https://chem.echa.europa.eu/activity-lists/svhcIdentification.
[25] European Chemicals Agency (ECHA). Authorisation Process List; 2025. https://chem.echa.europa.eu/activity-lists/authorisationProcess.
[26] European Chemicals Agency (ECHA). Dossier Evaluation List; 2025. https://chem.echa.europa.eu/activity-lists/dossierEvaluation.
[27] European Chemicals Agency (ECHA). CLH Process List; 2025. https://chem.echa.europa.eu/activity-lists/clhProcess.
[28] European Chemicals Agency (ECHA). Substance Evaluation List; 2025. https://chem.echa.europa.eu/activity-lists/substanceEvaluation.
[29] European Chemicals Agency (ECHA). POPs Process List; 2025. https://chem.echa.europa.eu/activity-lists/popsProcess.
[30] European Chemicals Agency (ECHA). PBT Assessment List; 2025. https://chem.echa.europa.eu/activity-lists/pbtAssessment.
[31] European Chemicals Agency (ECHA). Endocrine Disruptor Assessment List; 2025. https://chem.echa.europa.eu/activity-lists/edAssessment.
[32] European Chemicals Agency (ECHA). REACH Registrations Database; 2025. https://chem.echa.europa.eu/.
[33] European Commission. EU Pesticides Database: Active Substances; 2025. https://ec.europa.eu/food/plant/pesticides/eu-pesticides-database/start/screen/active-substances.
[34] OSPAR Commission. OSPAR List of Chemicals for Priority Action; 2024. https://www.ospar.org/work-areas/hasec/hazardous-substances/priority-action.
[35] Vlaamse Overheid. Sommatie Stoffenlijst (Waterkwaliteitsmodel); 2025. https://datasets.omgeving.vlaanderen.be/be.vlaanderen.omgeving.data.id.distribution.codelijst-sommatie_stoffen.3.0.8.sommatie_stoffen.csv.
[36] Rousseeuw PJ. Silhouettes: A Graphical Aid to the Interpretation and Validation of Cluster Analysis. Journal of Computational and Applied Mathematics. 1987;20:53-65.
[37] Lloyd SP. Least Squares Quantization in PCM. IEEE Transactions on Information Theory. 1982;28(2):129-37.
[38] McInnes L, Healy J, Melville J. UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction. arXiv preprint arXiv:180203426. 2018. Available from: https://arxiv.org/abs/1802.03426.
[39] Manning CD, Raghavan P, Schütze H. Introduction to Information Retrieval. Cambridge: Cambridge University Press; 2008.
[40] Wang W, Wei F, Dong L, Bao H, Yang N, Zhou M. MiniLM: Deep Self-Attention Distillation for Task-Agnostic Compression of Pre-Trained Transformers. In: Advances in Neural Information Processing Systems. vol. 33; 2020. p. 5776-88. Available from: https://proceedings.neurips.cc/paper/2020/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html.
[41] Devlin J, Chang MW, Lee K, Toutanova K. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In: Burstein J, Doran C, Solorio T, editors. Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers). Minneapolis, Minnesota: Association for Computational Linguistics; 2019. p. 4171-86. Available from: https://aclanthology.org/N19-1423/.
[42] Miles A, Bechhofer S. SKOS Simple Knowledge Organization System Reference. World Wide Web Consortium; 2009. Available from: https://www.w3.org/TR/skos-reference/.
[43] Sanderson R, Ciccarese P, Young B. Web Annotation Data Model. World Wide Web Consortium; 2017. Available from: https://www.w3.org/TR/annotation-model/.
[44] Reif DM, Martin MT, Tan SW, Houck KA, Judson RS, Richard AM, et al. Endocrine Profiling and Prioritization of Environmental Chemicals Using ToxCast Data. Environmental Health Perspectives.

2010;118(12):1714-20.